

Ocenjevanje

homografije

Zbirke slik

Predpriprava MS Coco 2017 dataset

```
from Generator import prepair_dataset

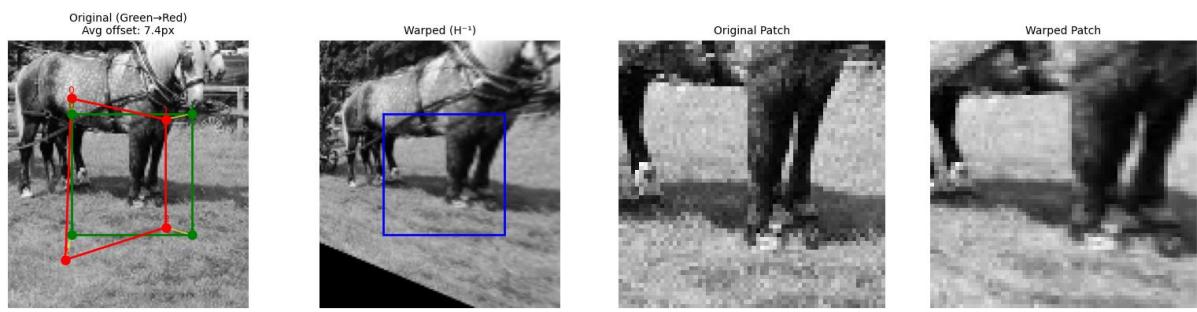
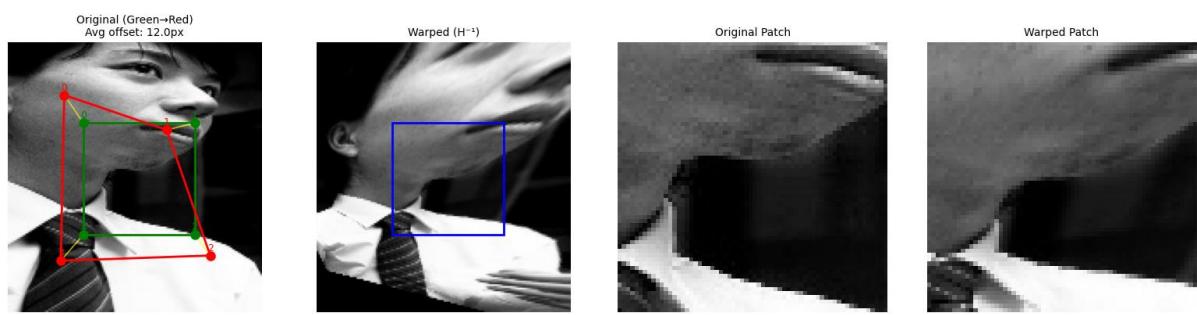
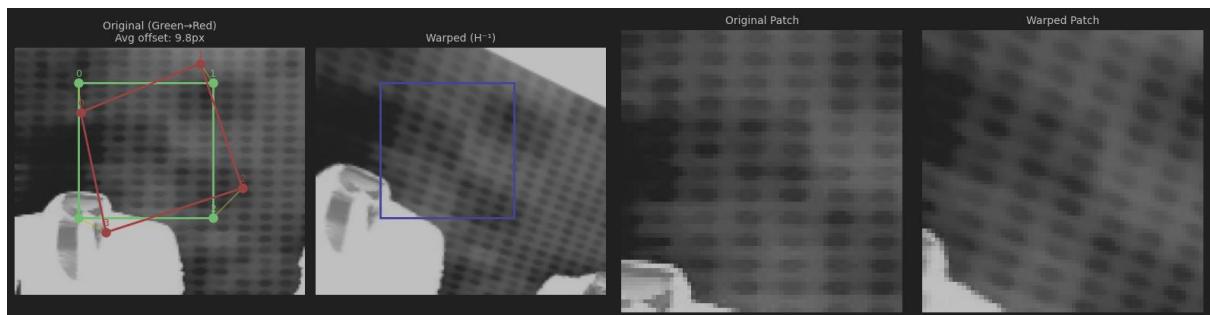
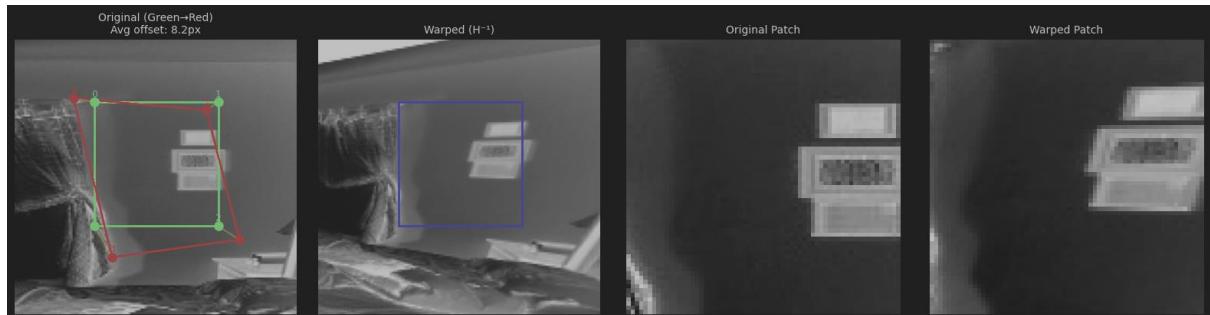
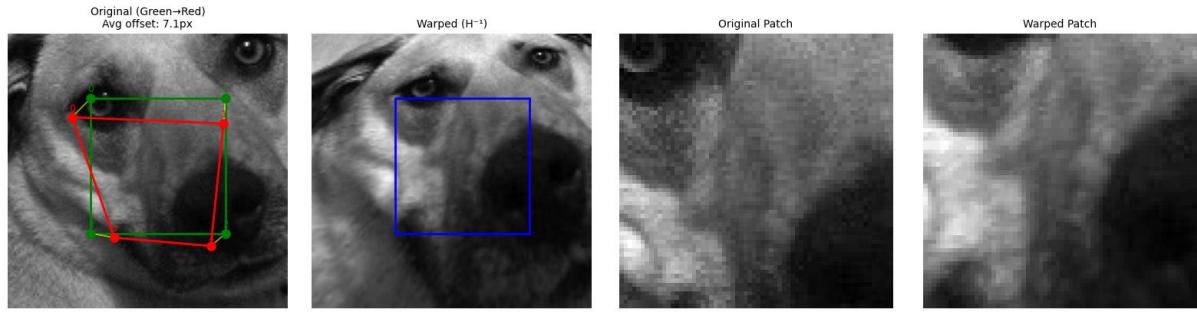
INPUT_DIR = "datasets/val2017"
PREPROCESSED_DIR = "datasets/val2017_preprocessed"

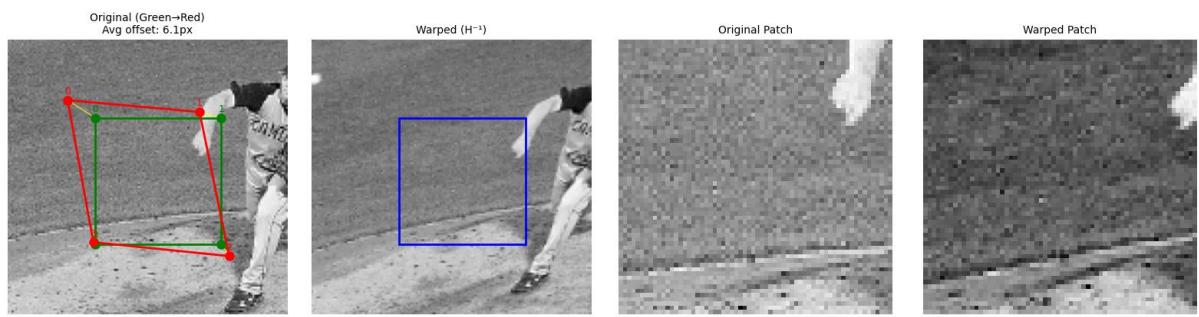
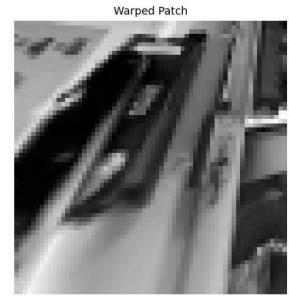
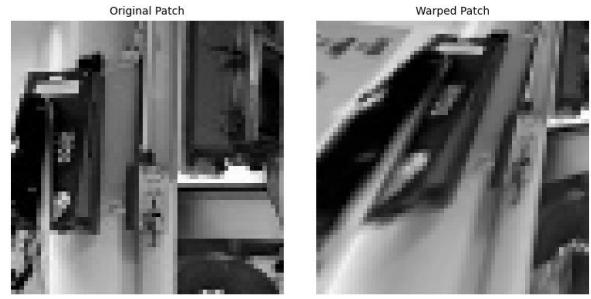
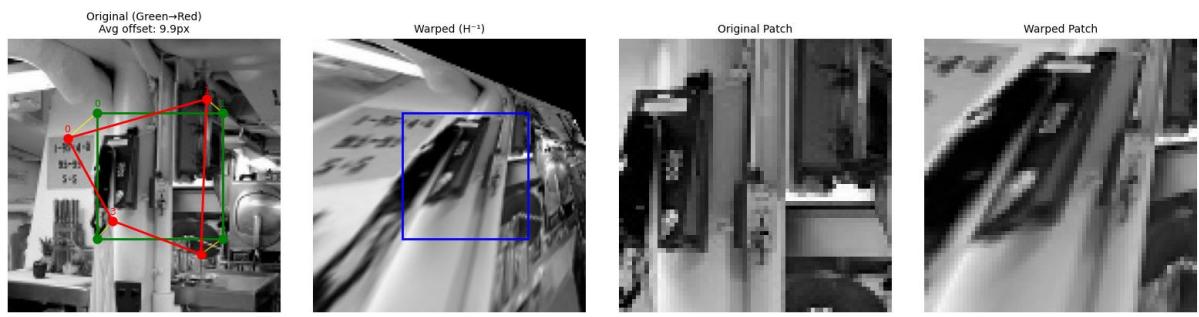
TARGET_SIZE = (320, 240)
prepair_dataset(INPUT_DIR, PREPROCESSED_DIR, TARGET_SIZE)
```

Generirane testnih primerov

```
from Generator import visualize_generate_pair

# Run visualization
for _ in range(3):
    visualize_generate_pair(PREPROCESSED_DIR)
```





Nevronski mreži

ResNet Block

```
from torchinfo import summary
from Models import ResNetBlock

# For ResNet block
block = ResNetBlock(in_channels=64, out_channels=128, stride=2,
dropout_rate=0.1)
summary(block, input_size=(1, 64, 32, 32))
```

Layer (type:depth-idx)	Output Shape	Param #
------------------------	--------------	---------

ResNetBlock	[1, 128, 16, 16]	--
─Conv2d: 1-1	[1, 128, 16, 16]	8,320
─Conv2d: 1-2	[1, 128, 16, 16]	73,856
─Dropout2d: 1-3	[1, 128, 16, 16]	--
─ReLU: 1-4	[1, 128, 16, 16]	--
─Conv2d: 1-5	[1, 128, 16, 16]	147,584
─Dropout2d: 1-6	[1, 128, 16, 16]	--
─ReLU: 1-7	[1, 128, 16, 16]	--

Total params: 229,760

Trainable params: 229,760

Non-trainable params: 0

Total mult-adds (Units.MEGABYTES): 58.82

Input size (MB): 0.26

Forward/backward pass size (MB): 0.79

Params size (MB): 0.92

Estimated Total Size (MB): 1.97

ResNet Body

```
from torchinfo import summary
from Models import ResNetBody

# For ResNet body
body = ResNetBody(in_channels=2, dropout_rate=0.1)
summary(body, input_size=(1, 2, 64, 64), depth=2)
```

Layer (type:depth-idx)	Output Shape	Param #
<hr/>		
ResNetBody	[1, 512]	--
└ Sequential: 1-1	[1, 64, 32, 32]	--
└ ResNetBlock: 2-1	[1, 64, 64, 64]	38,336
└ ResNetBlock: 2-2	[1, 64, 64, 64]	73,856
└ BatchNorm2d: 2-3	[1, 64, 64, 64]	128
└ ReLU: 2-4	[1, 64, 64, 64]	--
└ MaxPool2d: 2-5	[1, 64, 32, 32]	--
└ Sequential: 1-2	[1, 64, 16, 16]	--
└ ResNetBlock: 2-6	[1, 64, 32, 32]	73,856
└ ResNetBlock: 2-7	[1, 64, 32, 32]	73,856
└ BatchNorm2d: 2-8	[1, 64, 32, 32]	128
└ ReLU: 2-9	[1, 64, 32, 32]	--
└ MaxPool2d: 2-10	[1, 64, 16, 16]	--
└ Sequential: 1-3	[1, 128, 8, 8]	--
└ ResNetBlock: 2-11	[1, 128, 16, 16]	229,760
└ ResNetBlock: 2-12	[1, 128, 16, 16]	295,168
└ BatchNorm2d: 2-13	[1, 128, 16, 16]	256
└ ReLU: 2-14	[1, 128, 16, 16]	--
└ MaxPool2d: 2-15	[1, 128, 8, 8]	--
└ Sequential: 1-4	[1, 128, 8, 8]	--
└ ResNetBlock: 2-16	[1, 128, 8, 8]	295,168
└ ResNetBlock: 2-17	[1, 128, 8, 8]	295,168
└ BatchNorm2d: 2-18	[1, 128, 8, 8]	256
└ ReLU: 2-19	[1, 128, 8, 8]	--
└ Flatten: 1-5	[1, 8192]	--
└ Linear: 1-6	[1, 512]	4,194,816
<hr/>		

Total params: 5,570,752

Trainable params: 5,570,752

Non-trainable params: 0

Total mult-adds (Units.MEGABYTES): 787.15

=====

Input size (MB): 0.03

Forward/backward pass size (MB): 17.11

Params size (MB): 22.28

Estimated Total Size (MB): 39.42

=====

Homography Regressor

```
from torchinfo import summary
```

```
# For regression model
model = HomographyRegressor()
summary(model, input_size=(1, 2, 64, 64))
```

```
=====
Layer (type:depth-idx)          Output Shape       Param #
=====
HomographyRegressor            [1, 8]           --
|--ResNetBody: 1-1             [1, 512]          --
|  |--Sequential: 2-1          [1, 64, 32, 32]   112,320
|  |--Sequential: 2-2          [1, 64, 16, 16]   147,840
|  |--Sequential: 2-3          [1, 128, 8, 8]    525,184
|  |--Sequential: 2-4          [1, 128, 8, 8]    590,592
|  |--Flatten: 2-5             [1, 8192]         --
|  |--Linear: 2-6              [1, 512]          4,194,816
|  |--RegressionHead: 1-2     [1, 8]           --
|    |--Linear: 2-7             [1, 8]           4,104
=====
Total params: 5,574,856
Trainable params: 5,574,856
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 787.16
=====
Input size (MB): 0.03
Forward/backward pass size (MB): 17.11
Params size (MB): 22.30
Estimated Total Size (MB): 39.44
=====
```

Homography Classifier

```
from torchinfo import summary

# For classification model
model = HomographyClassifier(num_classes=21, class_dim=8)
summary(model, input_size=(1, 2, 64, 64), depth=2)
```

Layer (type:depth-idx)	Output Shape	Param #
<hr/>		
HomographyClassifier	[1, 8, 21]	--
--ResNetBody: 1-1	[1, 512]	--
└Sequential: 2-1	[1, 64, 32, 32]	112,320
└Sequential: 2-2	[1, 64, 16, 16]	147,840
└Sequential: 2-3	[1, 128, 8, 8]	525,184
└Sequential: 2-4	[1, 128, 8, 8]	590,592
└Flatten: 2-5	[1, 8192]	--
└Linear: 2-6	[1, 512]	4,194,816
--ClassificationHead: 1-2	[1, 8, 21]	--
└Linear: 2-7	[1, 168]	86,184
<hr/>		

Total params: 5,656,936

Trainable params: 5,656,936

Non-trainable params: 0

Total mult-adds (Units.MEGABYTES): 787.24

Input size (MB): 0.03

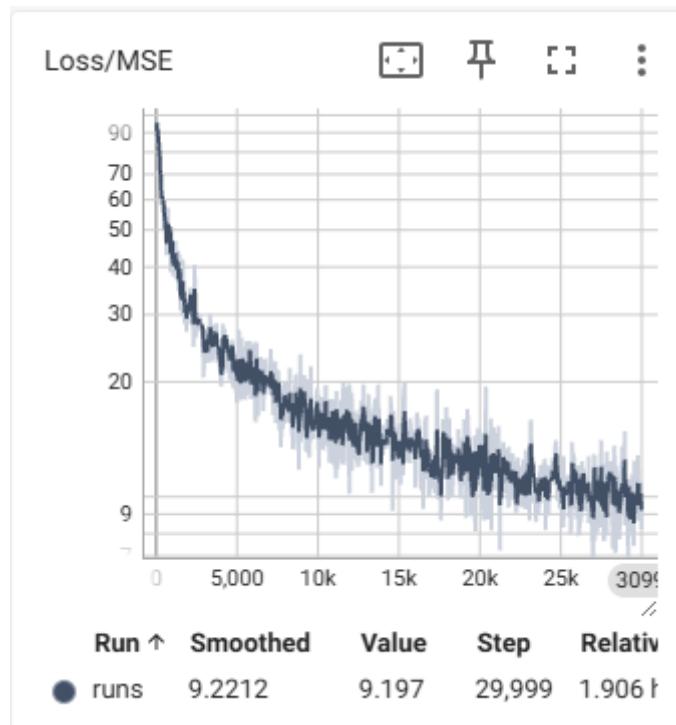
Forward/backward pass size (MB): 17.11

Params size (MB): 22.63

Estimated Total Size (MB): 39.77

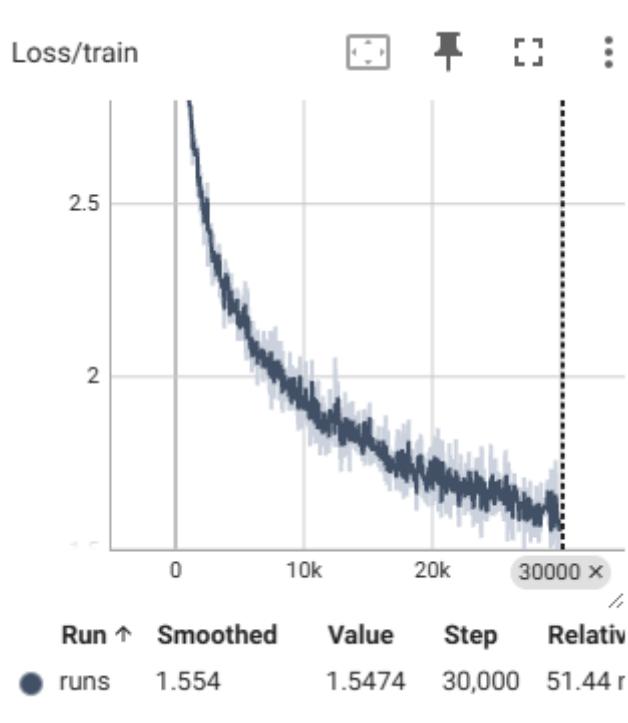
Učenje

Regression model (30k epochs, 64 samples)



(LOGARITEMSKA SKALA)

Classification model (30k epochs, 64 samples)



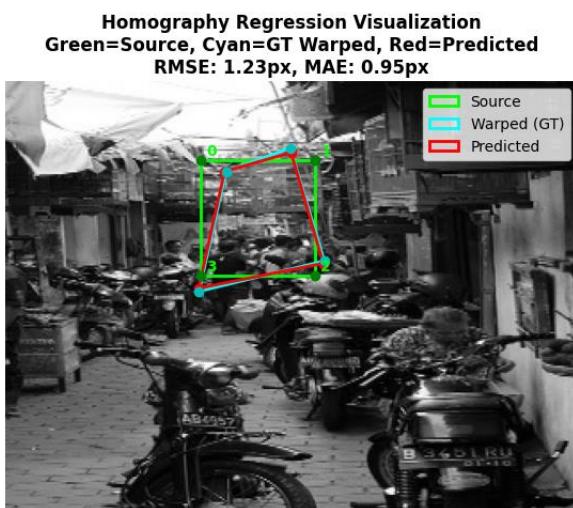
Rezultati

Regresija

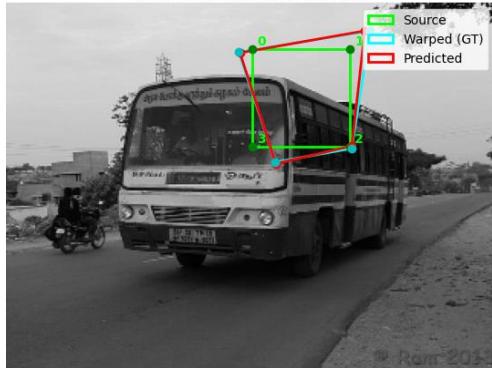
visualization of regression results

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = HomographyRegressor().to(device)
state = torch.load("checkpoints_homography_regressor_all/h_regressor_all.pth")
model.load_state_dict(state)
model.eval()

img = get_random_images(1, image_dir=PREPROCESSED_DIR)[0]
visualize_regression_result(model=model, image=img)
```



Homography Regression Visualization
Green=Source, Cyan=GT Warped, Red=Predicted
RMSE: 1.43px, MAE: 0.95px



Homography Regression Visualization
Green=Source, Cyan=GT Warped, Red=Predicted
RMSE: 1.67px, MAE: 1.58px

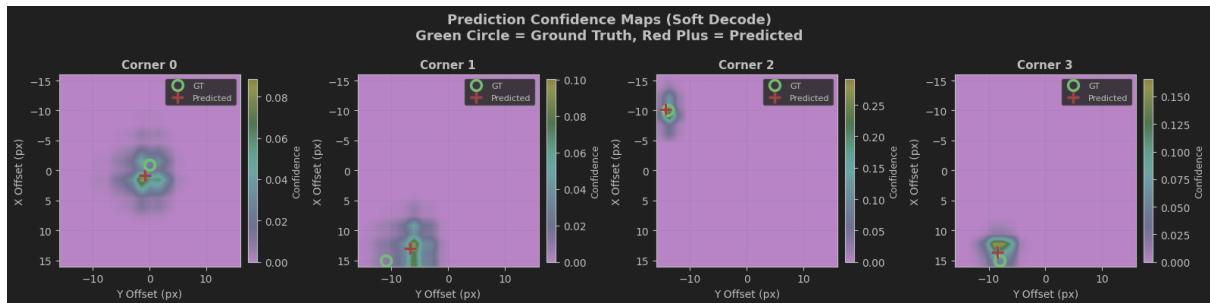
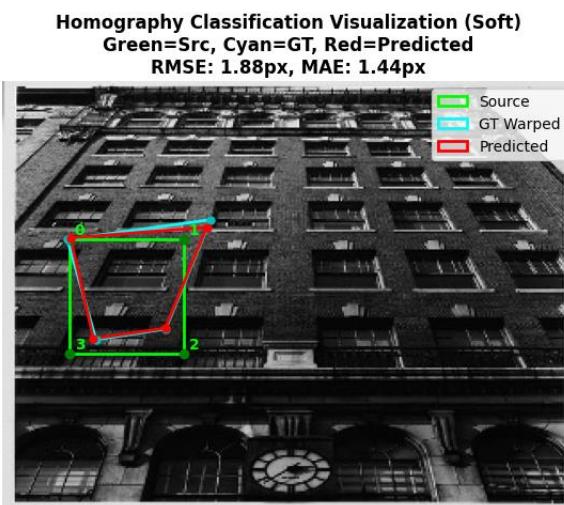


Klasifikasi

```
# visualization of classification results
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = HomographyClassifier(num_classes=21, class_dim=8).to(device)
state = torch.load("checkpoints_homography_classify_all/checkpoint_epoch_30000.pth")["model_state_dict"]
model.load_state_dict(state)
model.eval()

img = get_random_images(1, image_dir=PREPROCESSED_DIR)[0]
visualize_classification_result(model=model, image=img, soft_decode=True)
```



Homography Classification Visualization (Soft)

Green=Src, Cyan=GT, Red=Predicted

RMSE: 2.67px, MAE: 2.35px



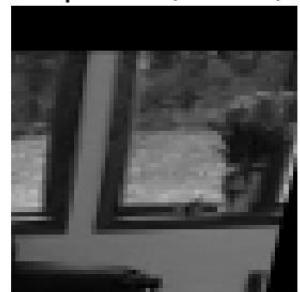
Original Patch



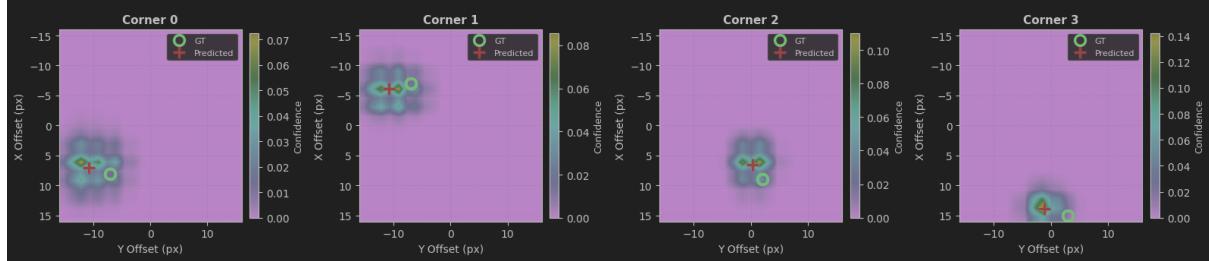
Warped Patch (GT)



Warped Patch (Predicted)



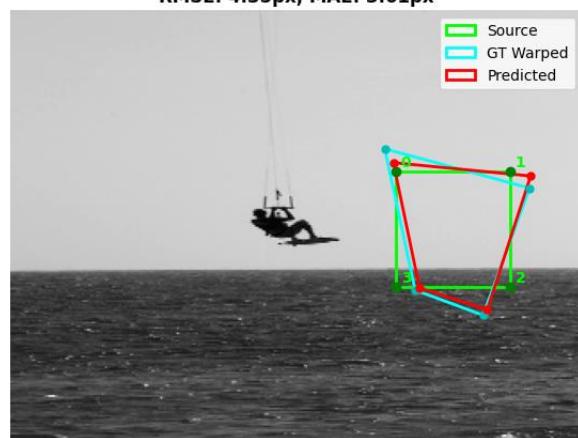
Prediction Confidence Maps (Soft Decode)
Green Circle = Ground Truth, Red Plus = Predicted



Homography Classification Visualization (Soft)

Green=Src, Cyan=GT, Red=Predicted

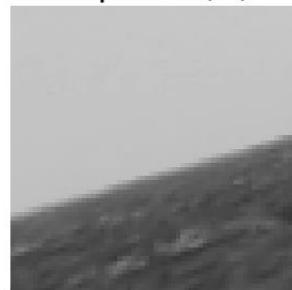
RMSE: 4.35px, MAE: 3.61px



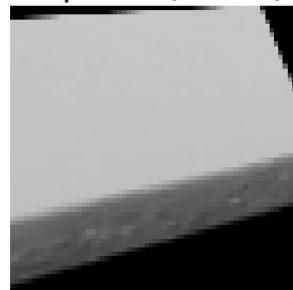
Original Patch



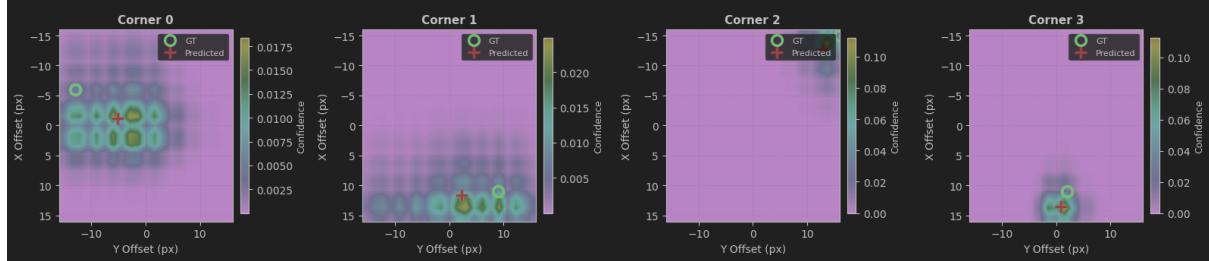
Warped Patch (GT)



Warped Patch (Predicted)



Prediction Confidence Maps (Soft Decode)
Green Circle = Ground Truth, Red Plus = Predicted



Evaluacija

Regression models

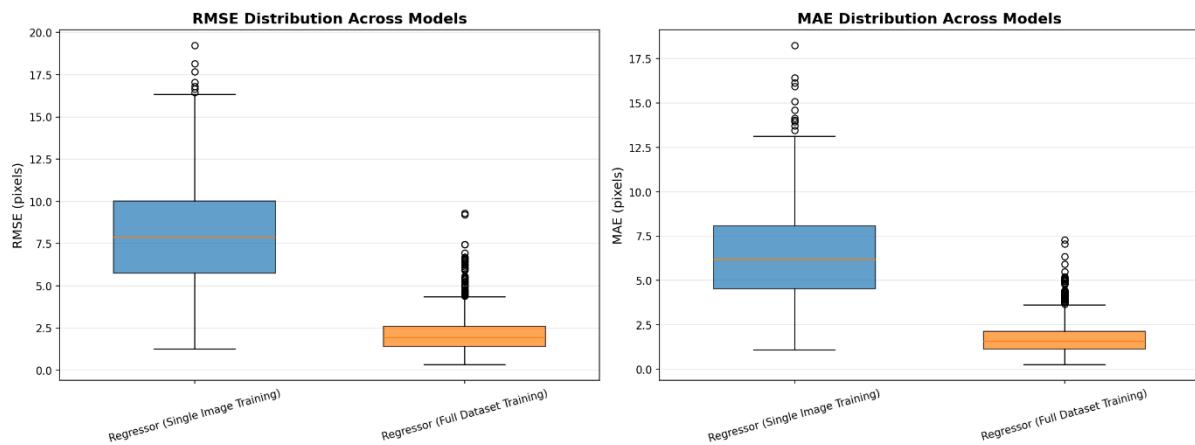
```
images = get_random_images(100, "datasets/val2017_preprocessed")
test_samples = generate_test_set(images, samples_per_image=10)

# Compare regression models
regressor_single = HomographyRegressor()
regressor_single.load_state_dict(torch.load("checkpoints_homography_regressor_single/h_regressor_single.pth"))

regressor_all = HomographyRegressor()
regressor_all.load_state_dict(torch.load("checkpoints_homography_regressor_all/h_regressor_all.pth"))

results = {
    'Regressor (Single Image Training)': evaluate_regressor(regressor_single, test_samples),
    'Regressor (Full Dataset Training)': evaluate_regressor(regressor_all, test_samples)
}

summarize_and_plot(results, save_dir="eval_results")
```



EVALUATION SUMMARY

ocena z Regressor (Single Image Training) mean: 10.29 std: 6.52

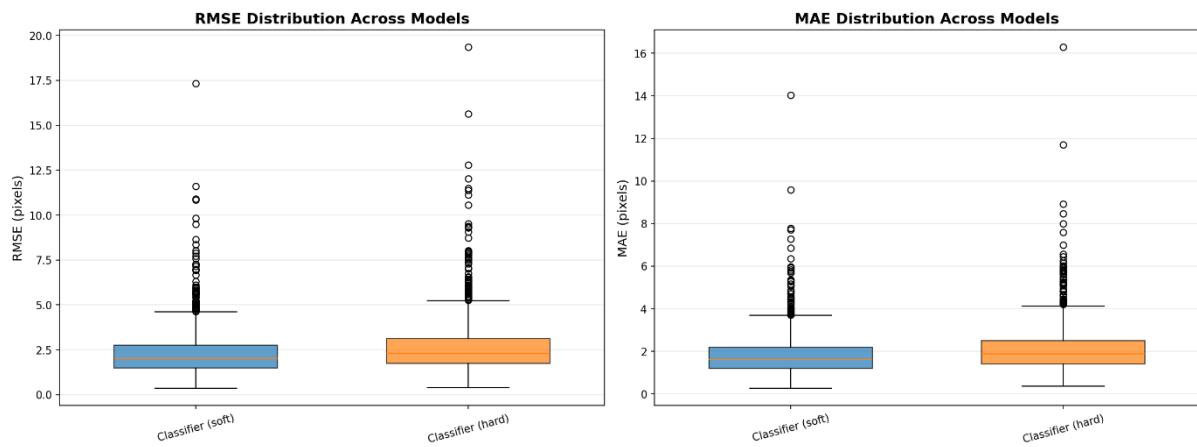
ocena z Regressor (Full Dataset Training) mean: 2.77 std: 2.11

Classification models

```
# Compare classifiers (soft vs hard decoding)
classifier = HomographyClassifier()
classifier.load_state_dict(torch.load("checkpoints_homography_classify_all/h_classify_all.pth"))

results = {
    'Classifier (soft)': evaluate_classifier(classifier, test_samples, soft_decode=True),
    'Classifier (hard)': evaluate_classifier(classifier, test_samples, soft_decode=False)
}

summarize_and_plot(results, save_dir="eval_results")
```



EVALUATION SUMMARY

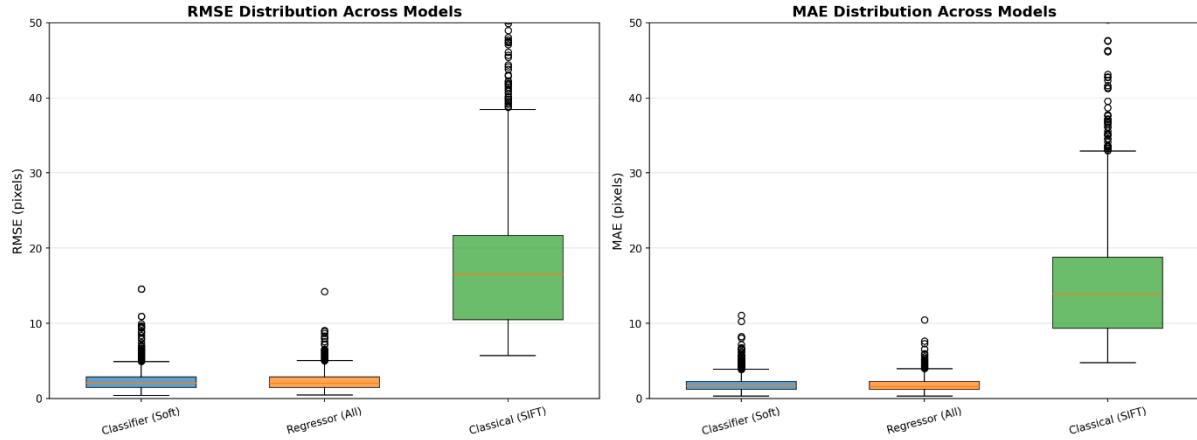
ocena z Classifier (soft) mean: 2.95 std: 2.56

ocena z Classifier (hard) mean: 3.40 std: 3.05

Comparing best models with classical

```
results = {
    'Classifier (Soft)': evaluate_classifier(classifier, test_samples, soft_decode=True),
    'Regressor (All)': evaluate_regressor(regressor_all, test_samples),
    'Classical (SIFT)': evaluate_classical(test_samples)
}

summarize_and_plot(results, save_dir="eval_results", ylim=50)
```



EVALUATION SUMMARY

ocena z Classifier (Soft) mean: 3.05 std: 2.74

ocena z Regressor (All) mean: 2.96 std: 2.37

ocena z Classical (SIFT) mean: 24.30 std: 21.01